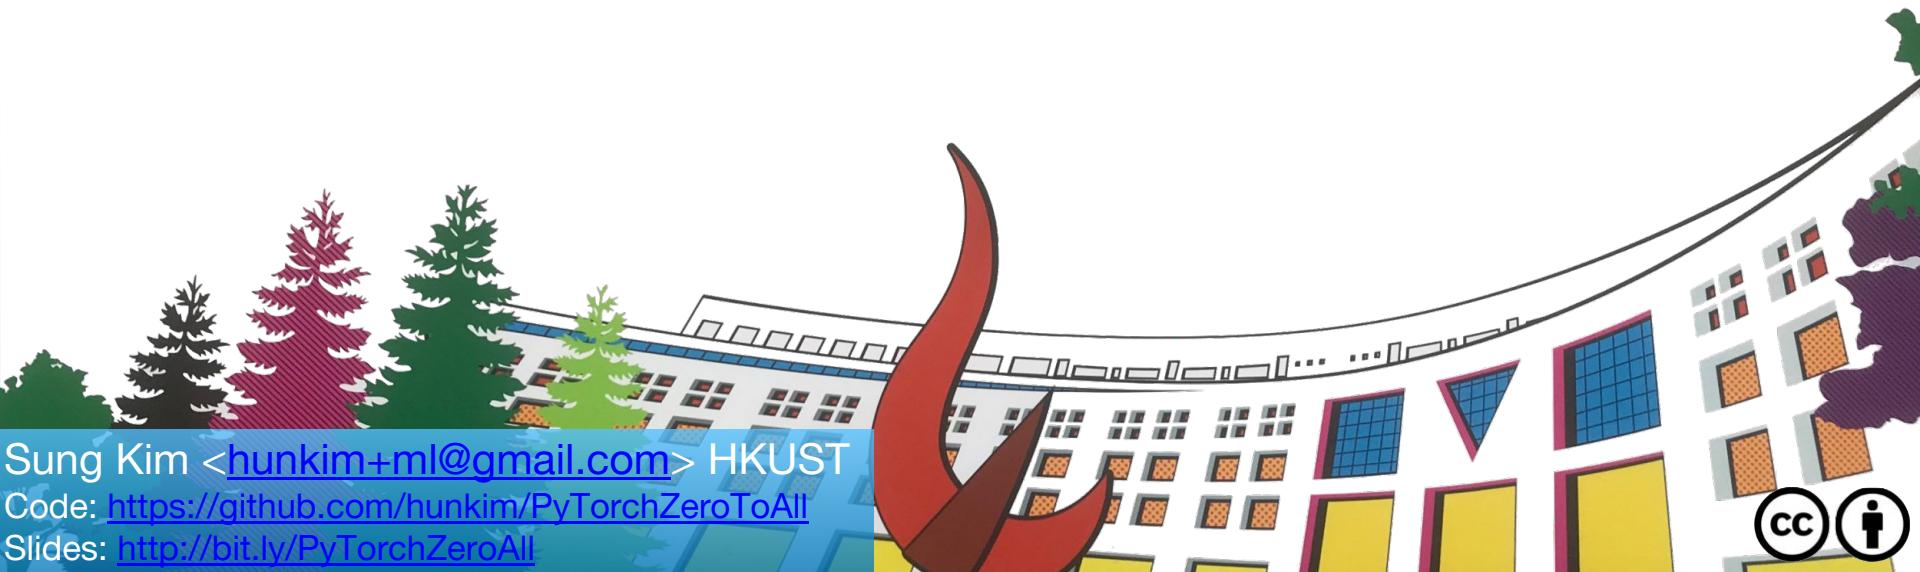


ML/DL for Everyone with PYTORCH

Lecture 13: RNN II



Sung Kim <hunkim+ml@gmail.com> HKUST
Code: <https://github.com/hunkim/PyTorchZeroToAll>
Slides: <http://bit.ly/PyTorchZeroAll>



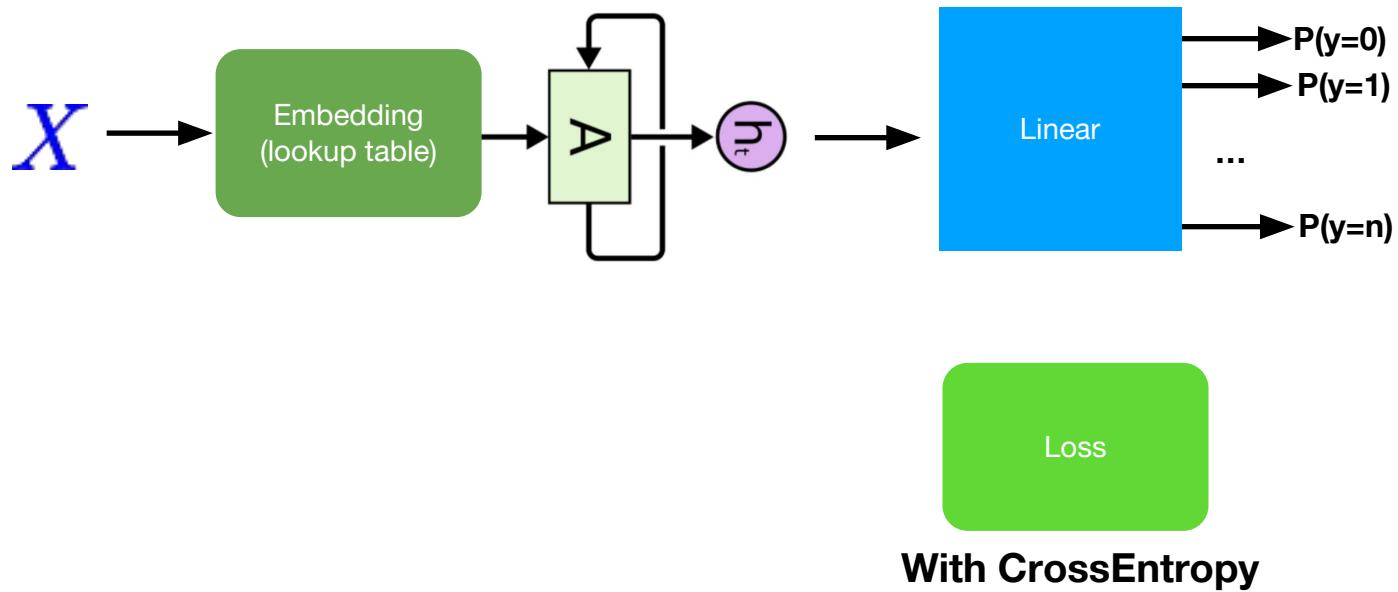
Call for Comments

Please feel free to add comments directly on these slides.

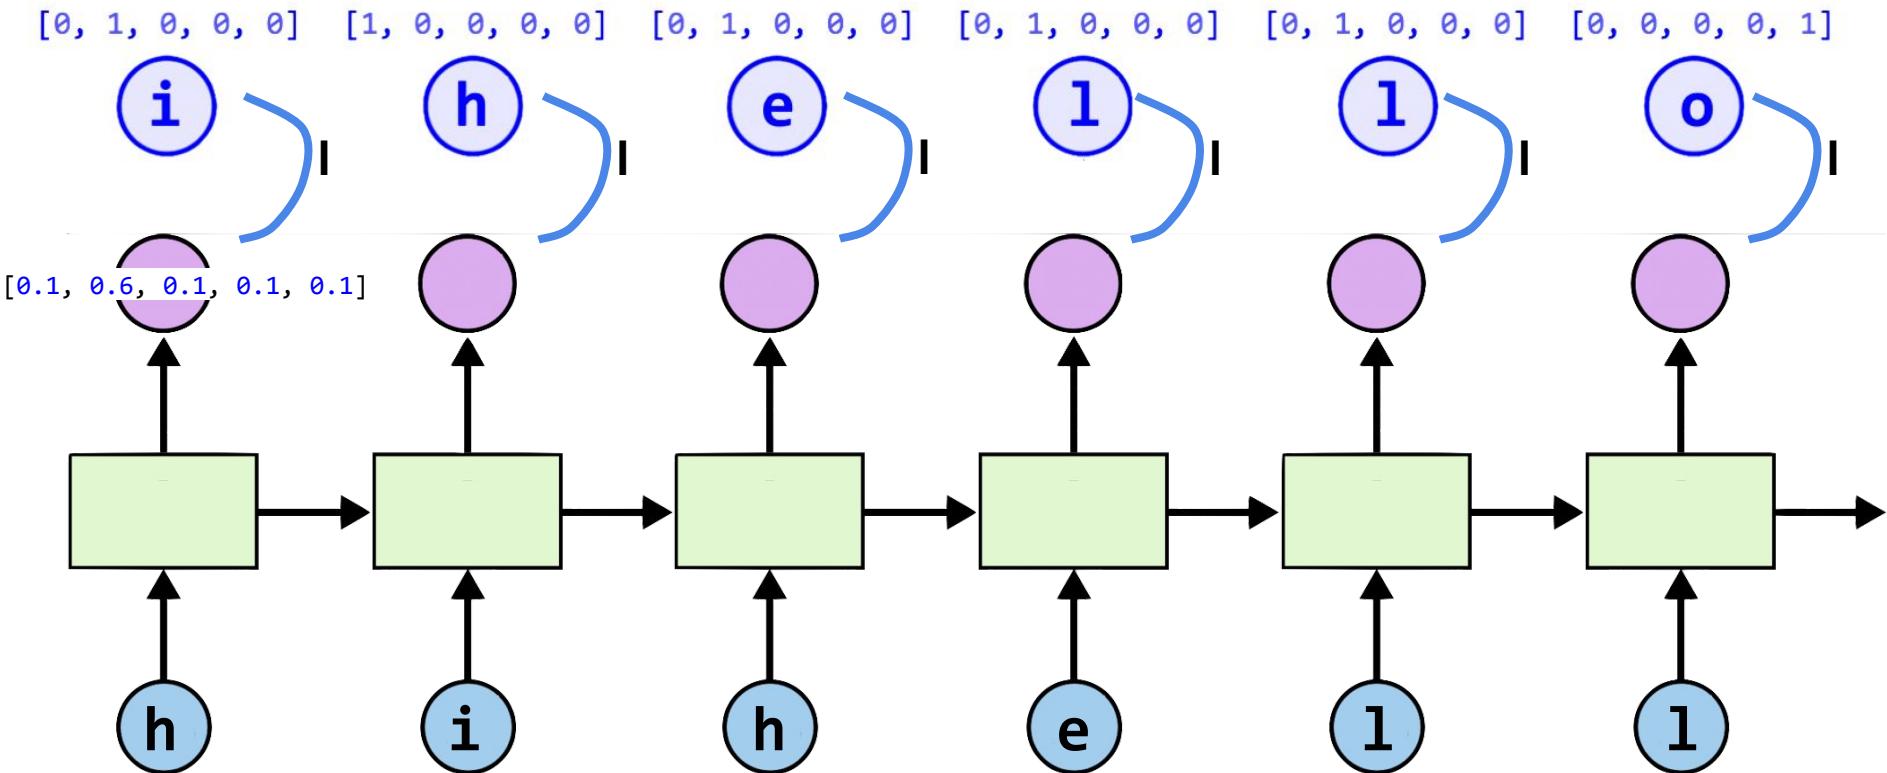
Other slides: <http://bit.ly/PyTorchZeroAll>



Typical RNN Models

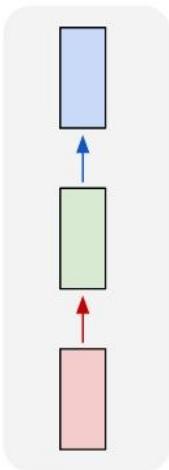


RNN Loss and training

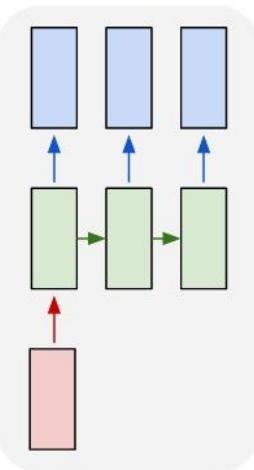


RNN Applications

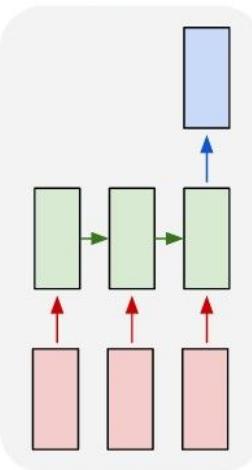
one to one



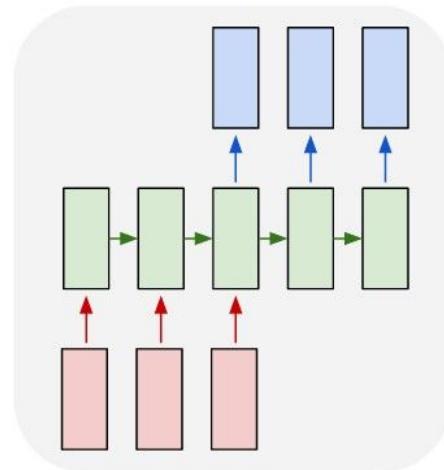
one to many



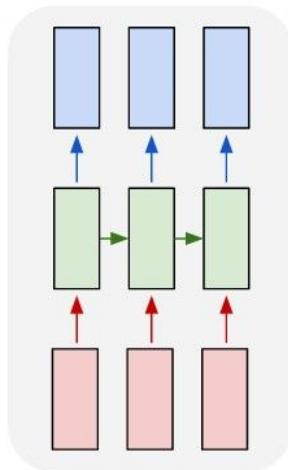
many to one



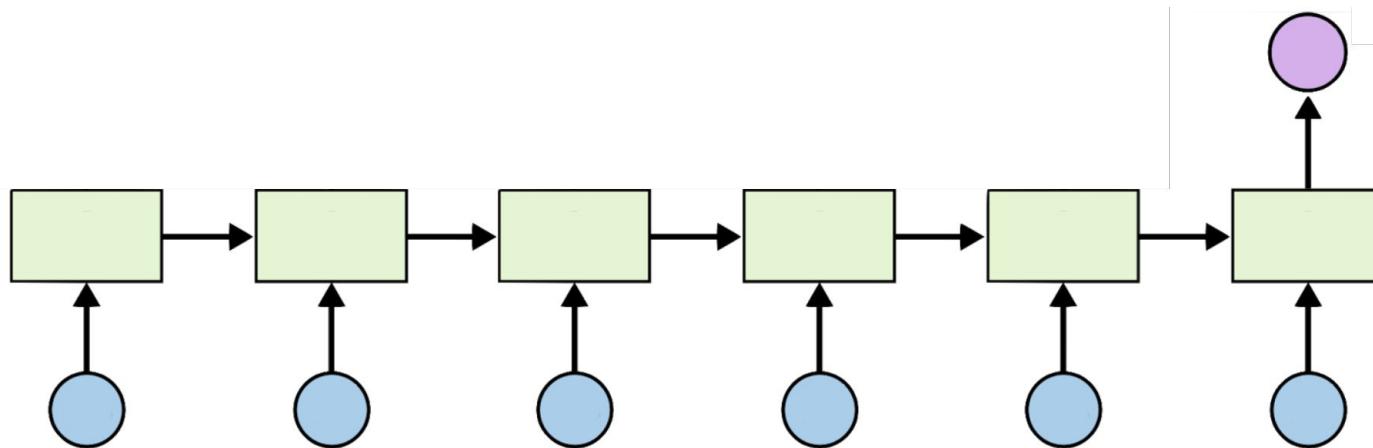
many to many



many to many



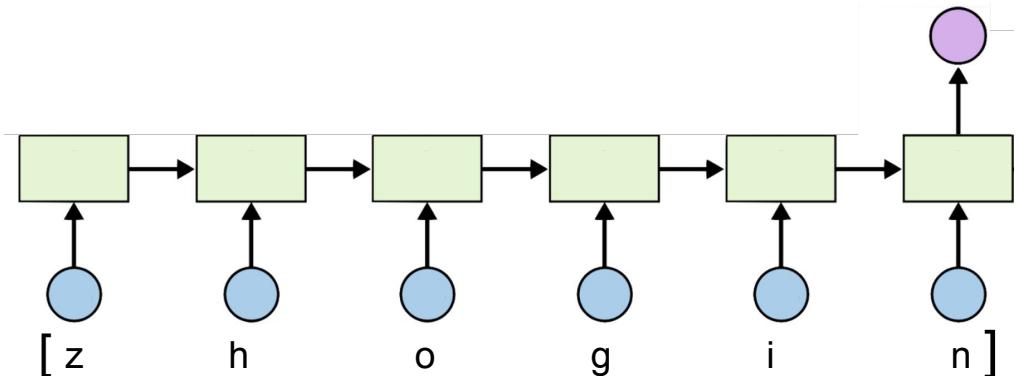
RNN Classification



Name Classification

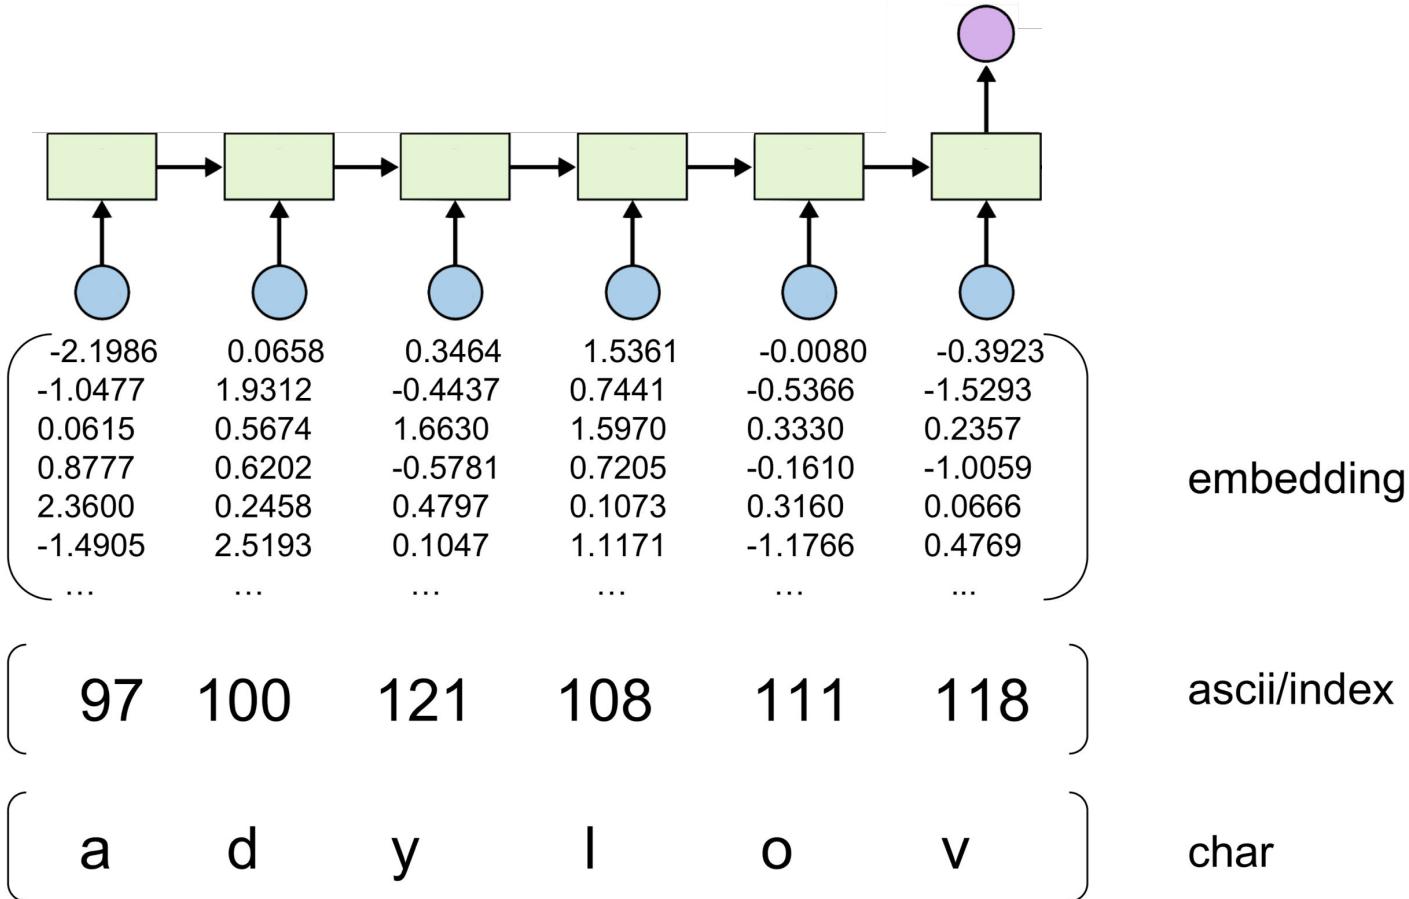
0	Nader	Arabic
1	Malouf	Arabic
2	Terajima	Japanese
3	Huie	Chinese
4	Chertushkin	Russian
5	Davletkildeev	Russian
6	Movchun	Russian
7	Pokhvoschev	Russian
8	Zhogin	Russian
9	Hancock	English
10	Tomkins	English

Softmax output
(18 countries)

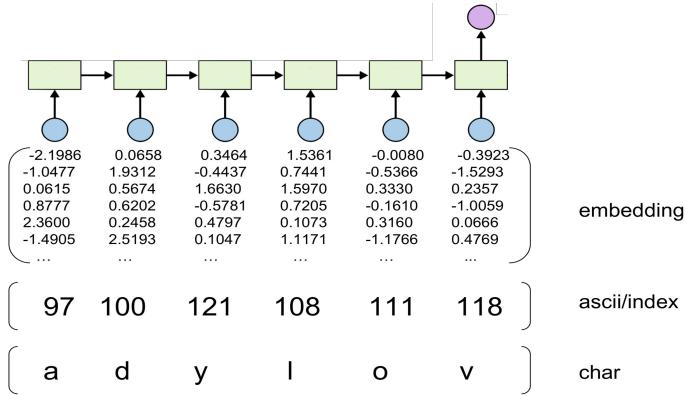


Input representations

Softmax output
(18 countries)



Implementation



```
if __name__ == '__main__':
    classifier = RNNClassifier(N_CHARS,
                                HIDDEN_SIZE, N_CLASSES)

    arr, _ = str2ascii_arr('adylov')
    inp = Variable(torch.LongTensor([arr]))
    out = classifier(inp)
    print("in", inp.size(), "out", out.size())
    # in torch.Size([1, 6])
    #out torch.Size([1, 1, 18])
```

```
class RNNClassifier(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, n_layers=1):
        super(RNNClassifier, self).__init__()
        self.hidden_size = hidden_size
        self.n_layers = n_layers
        self.embedding = nn.Embedding(input_size, hidden_size)
        self.gru = nn.GRU(hidden_size, hidden_size, n_layers)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, input):
        # Note: we run this all at once (over the whole input sequence)
        # input = B x S . size(0) = B
        batch_size = input.size(0)

        # input: B x S -- (transpose) --> S x B
        input = input.t()

        # Embedding S x B -> S x B x I (embedding size)
        print(" input", input.size())
        embedded = self.embedding(input)
        print(" embedding", embedded.size())

        # Make a hidden
        hidden = self._init_hidden(batch_size)
        output, hidden = self.gru(embedded, hidden)
        print(" gru hidden output", hidden.size())
        # Use the last layer output as FC's input
        # No need to unpack, since we are going to use hidden
        fc_output = self.fc(hidden)
        print(" fc output", fc_output.size())
        return fc_output

    def _init_hidden(self, batch_size):
        hidden = torch.zeros(self.n_layers, batch_size, self.hidden_size)
        return Variable(hidden)
```

Implementation

[97	100	121	108	111	118]
transpose	97	-2.1986	-1.0477	0.0615	0.8777	2.3600	-1.4905
100	0.0658	1.9312	0.5674	0.6202	0.2458	2.5193	
121	0.3464	-0.4437	1.6630	-0.5781	0.4797	0.1047	
108	1.5361	0.7441	1.5970	0.7205	0.1073	1.1171	
111	-0.0080	-0.5366	0.3330	-0.1610	0.3160	-1.1766	
118	-0.3923	-1.5293	0.2357	-1.0059	0.0666	0.4769	
adylov							Embedding

```
if __name__ == '__main__':
    classifier = RNNClassifier(N_CHARS,
                                HIDDEN_SIZE, N_CLASSES)

    arr, _ = str2ascii_arr('adylov')
    inp = Variable(torch.LongTensor([arr]))
    out = classifier(inp)
    print("in", inp.size(), "out", out.size())
    # in torch.Size([1, 6])
    #out torch.Size([1, 1, 18])
```

```
class RNNClassifier(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, n_layers=1):
        super(RNNClassifier, self).__init__()
        self.hidden_size = hidden_size
        self.n_layers = n_layers
        self.embedding = nn.Embedding(input_size, hidden_size)
        self.gru = nn.GRU(hidden_size, hidden_size, n_layers)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, input):
        # Note: we run this all at once (over the whole input sequence)
        # input = B x S . size(0) = B
        batch_size = input.size(0)

        # input: B x S -- (transpose) --> S x B
        input = input.t()

        # Embedding S x B -> S x B x I (embedding size)
        print(" input", input.size())
        embedded = self.embedding(input)
        print(" embedding", embedded.size())

        # Make a hidden
        hidden = self._init_hidden(batch_size)
        output, hidden = self.gru(embedded, hidden)
        print(" gru hidden output", hidden.size())
        # Use the last layer output as FC's input
        # No need to unpack, since we are going to use hidden
        fc_output = self.fc(hidden)
        print(" fc output", fc_output.size())
        return fc_output

    def _init_hidden(self, batch_size):
        hidden = torch.zeros(self.n_layers, batch_size, self.hidden_size)
        return Variable(hidden)
```

Batch?

	[97	100	121	108	111	118]
transpose	→	97	-2.1986	-1.0477	0.0615	0.8777	2.3600	-1.4905
	100	0.0658	1.9312	0.5674	0.6202	0.2458	2.5193	
	121	0.3464	-0.4437	1.6630	-0.5781	0.4797	0.1047	
	108	1.5361	0.7441	1.5970	0.7205	0.1073	1.1171	
	111	-0.0080	-0.5366	0.3330	-0.1610	0.3160	-1.1766	
	118	-0.3923	-1.5293	0.2357	-1.0059	0.0666	0.4769	
adylov								Embedding

```
if __name__ == '__main__':
    classifier = RNNClassifier(N_CHARS,
                                HIDDEN_SIZE, N_CLASSES)

    arr, _ = str2ascii_arr('adylov')
    inp = Variable(torch.LongTensor([arr]))
    out = classifier(inp)
    print("in", inp.size(), "out", out.size())
    # in torch.Size([1, 6])
    # out torch.Size([1, 1, 18])
```

```
if __name__ == '__main__':
    names = ['adylov', 'solan', 'hard', 'san']
    classifier = RNNClassifier(N_CHARS, HIDDEN_SIZE, N_CLASSES)

    for name in names:
        arr, _ = str2ascii_arr(name)
        inp = Variable(torch.LongTensor([arr]))
        out = classifier(inp)
        print("in", inp.size(), "out", out.size())

    # in torch.Size([1, 6]) out torch.Size([1, 1, 18])
    # in torch.Size([1, 5]) out torch.Size([1, 1, 18])
    # ...
```

Batch?

```
[ 97 100 121 108 111 118 ]
```

transpose

```
97   -2.1986  -1.0477  0.0615  0.8777  2.3600  -1.4905  
100  0.0658  1.9312  0.5674  0.6202  0.2458  2.5193  
121  0.3464  -0.4437  1.6630  -0.5781  0.4797  0.1047  
108  1.5361  0.7441  1.5970  0.7205  0.1073  1.1171  
111  -0.0080  -0.5366  0.3330  -0.1610  0.3160  -1.1766  
118  -0.3923  -1.5293  0.2357  -1.0059  0.0666  0.4769
```

adylov

Embedding

```
if __name__ == '__main__':  
    classifier = RNNClassifier(N_CHARS,  
                                HIDDEN_SIZE, N_CLASSES)  
  
    arr, _ = str2ascii_arr('adylov')  
    inp = Variable(torch.LongTensor([arr]))  
    out = classifier(inp)  
    print("in", inp.size(), "out", out.size())  
    # in torch.Size([1, 6]) out torch.Size([1, 1, 18])  
    #out torch.Size([1, 1, 18])  
    # ...
```

adylov	sloan	harb	san
97	115	104	115
100	111	97	97
121	108	114	110
108	97	100	
111	110		
118			

```
if __name__ == '__main__':  
    names = ['adylov', 'solan', 'hard', 'san']  
    classifier = RNNClassifier(N_CHARS, HIDDEN_SIZE, N_CLASSES)  
  
    for name in names:  
        arr, _ = str2ascii_arr(name)  
        inp = Variable(torch.LongTensor([arr]))  
        out = classifier(inp)  
        print("in", inp.size(), "out", out.size())  
  
        # in torch.Size([1, 6]) out torch.Size([1, 1, 18])  
        # in torch.Size([1, 5]) out torch.Size([1, 1, 18])  
        # ...
```

Zero padding

adylov sloan harb san

97	115	104	115
100	111	97	97
121	108	114	110
108	97	100	
111	110		
118			

Zero padding

adylov	sloan	harb	san
97	115	104	115
100	111	97	97
121	108	114	110
108	97	100	
111	110		
118			



adylov	sloan	harb	san
97	115	104	115
100	111	97	97
121	108	114	110
108	97	100	0
111	110	0	0
118	0	0	0

Zero padding

adylov	sloan	harb	san
97	115	104	115
100	111	97	97
121	108	114	110
108	97	100	
111	110		
118			



adylov	sloan	harb	san
97	115	104	115
100	111	97	97
121	108	114	110
108	97	100	0
111	110	0	0
118	0	0	0

```
def pad_sequences(vectorized_seqs, seq_lengths):
    seq_tensor = torch.zeros((len(vectorized_seqs), seq_lengths.max())).long()
    for idx, (seq, seq_len) in enumerate(zip(vectorized_seqs, seq_lengths)):
        seq_tensor[idx, :seq_len] = torch.LongTensor(seq)
    return seq_tensor
```

Embedding

adylov	sloan	harb	san
97	115	104	115
100	111	97	97
121	108	114	110
108	97	100	0
111	110	0	0
118	0	0	0



$\begin{pmatrix} -2.1986 & -1.0477 & 0.0615 & 0.8777 & 2.3600 & -1.4905 \\ 0.4186 & 0.4627 & -1.0357 & -0.9848 & 3.0676 & 0.2172 \\ -1.8650 & -0.6240 & -1.7311 & 0.7280 & 0.3623 & -0.1245 \\ 0.4186 & 0.4627 & -1.0357 & -0.9848 & 3.0676 & 0.2172 \end{pmatrix}$	A	S	H	S
$\begin{pmatrix} 0.0658 & 1.9312 & 0.5674 & 0.6202 & 0.2458 & 2.5193 \\ 1.5361 & 0.7441 & 1.5970 & 0.7205 & 0.1073 & 1.1171 \\ 1.5617 & -0.6109 & 1.3775 & 0.7142 & 1.2291 & 0.5016 \\ 1.5617 & -0.6109 & 1.3775 & 0.7142 & 1.2291 & 0.5016 \end{pmatrix}$	d	I	a	a
$\begin{pmatrix} 0.3464 & -0.4437 & 1.6630 & -0.5781 & 0.4797 & 0.1047 \\ -0.0080 & -0.5366 & 0.3330 & -0.1610 & 0.3160 & -1.1766 \\ 1.5768 & -0.6135 & -0.6442 & 0.1905 & 0.3791 & 0.1173 \\ 0.3804 & -1.9142 & 1.2423 & 1.1947 & -0.7331 & -1.2344 \end{pmatrix}$	y	o	r	n
$\begin{pmatrix} 1.5361 & 0.7441 & 1.5970 & 0.7205 & 0.1073 & 1.1171 \\ 1.5617 & -0.6109 & 1.3775 & 0.7142 & 1.2291 & 0.5016 \\ 0.6582 & 1.7296 & -0.7542 & -0.6202 & 0.2330 & 1.7715 \\ 1.6387 & -1.2216 & -0.1584 & 1.9363 & -0.3439 & -0.2585 \end{pmatrix}$	I	a	b	ø
$\begin{pmatrix} -0.0080 & -0.5366 & 0.3330 & -0.1610 & 0.3160 & -1.1766 \\ 0.3804 & -1.9142 & 1.2423 & 1.1947 & -0.7331 & -1.2344 \\ 1.6387 & -1.2216 & -0.1584 & 1.9363 & -0.3439 & -0.2585 \\ 1.6387 & -1.2216 & -0.1584 & 1.9363 & -0.3439 & -0.2585 \end{pmatrix}$	o	n	ø	ø
$\begin{pmatrix} -0.3923 & -1.5293 & 0.2357 & -1.0059 & 0.0666 & 0.4769 \\ 1.6387 & -1.2216 & -0.1584 & 1.9363 & -0.3439 & -0.2585 \\ 1.6387 & -1.2216 & -0.1584 & 1.9363 & -0.3439 & -0.2585 \\ 1.6387 & -1.2216 & -0.1584 & 1.9363 & -0.3439 & -0.2585 \end{pmatrix}$	v	ø	ø	ø

```
# Embedding S x B -> S x B x I (embedding size)
embedded = self.embedding(input)
```

Full implementation

```
def str2ascii_arr(msg):
    arr = [ord(c) for c in msg]
    return arr, len(arr)

# pad sequences and sort the tensor
def pad_sequences(vectorized_seqs, seq_lengths):
    seq_tensor = torch.zeros((len(vectorized_seqs), seq_lengths.max())).long()
    for idx, (seq, seq_len) in enumerate(zip(vectorized_seqs, seq_lengths)):
        seq_tensor[idx, :seq_len] = torch.LongTensor(seq)
    return seq_tensor

# Create necessary variables, lengths, and target
def make_variables(names):
    sequence_and_length = [str2ascii_arr(name) for name in names]
    vectorized_seqs = [sl[0] for sl in sequence_and_length]
    seq_lengths = torch.LongTensor([sl[1] for sl in sequence_and_length])
    return pad_sequences(vectorized_seqs, seq_lengths)

if __name__ == '__main__':
    names = ['adylov', 'solan', 'hard', 'san']
    classifier = RNNClassifier(N_CHARS, HIDDEN_SIZE, N_CLASSES)

    inputs = make_variables(names)
    out = classifier(inp)
    print("batch in", inp.size(), "batch out", out.size())

    # batch in torch.Size([4, 6]) batch out torch.Size([1, 4, 18])
```

Full implementation

```
def str2ascii_arr(msg):
    arr = [ord(c) for c in msg]
    return arr, len(arr)

# pad sequences and sort the tensor
def pad_sequences(vectorized_seqs, seq_lengths):
    seq_tensor = torch.zeros((len(vectorized_seqs), seq_lengths.max())).long()
    for idx, (seq, seq_len) in enumerate(zip(vectorized_seqs, seq_lengths)):
        seq_tensor[idx, :seq_len] = torch.LongTensor(seq)
    return seq_tensor

# Create necessary variables, lengths, and target
def make_variables(names):
    sequence_and_length = [str2ascii_arr(name) for name in names]
    vectorized_seqs = [sl[0] for sl in sequence_and_length]
    seq_lengths = torch.LongTensor([sl[1] for sl in sequence_and_length])
    return pad_sequences(vectorized_seqs, seq_lengths)

if __name__ == '__main__':
    names = ['adylov', 'solan', 'hard', 'san']
    classifier = RNNClassifier(N_CHARS, HIDDEN_SIZE, N_CLASSES)

    inputs = make_variables(names)
    out = classifier(inputs)
    print("batch in", inputs.size(), "batch out", out.size())

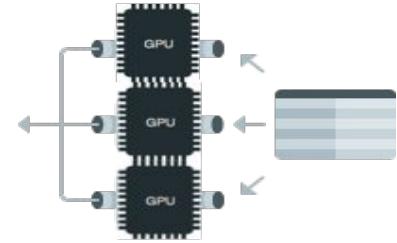
    # batch in torch.Size([4, 6]) batch out torch.Size([1, 4, 18])
```

```
optimizer =
torch.optim.Adam(classifier.parameters(),
                  lr=0.001)
criterion = nn.CrossEntropyLoss()

...
loss = criterion(output, target)

classifier.zero_grad()
loss.backward()
optimizer.step()
```

GPU/Data Parallel



1

Copy all variables to gpu

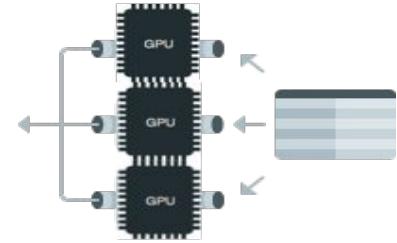
```
if torch.cuda.is_available():
    return Variable(tensor.cuda())
else:
    return Variable(tensor)
```

2

Put your models on gpu

```
classifier = RNNClassifier(N_CHARS, HIDDEN_SIZE, ...
if torch.cuda.is_available():
    classifier.cuda()
```

GPU/Data Parallel



1

Copy all variables to gpu

```
if torch.cuda.is_available():
    return Variable(tensor.cuda())
else:
    return Variable(tensor)
```

2

Put your models on gpu

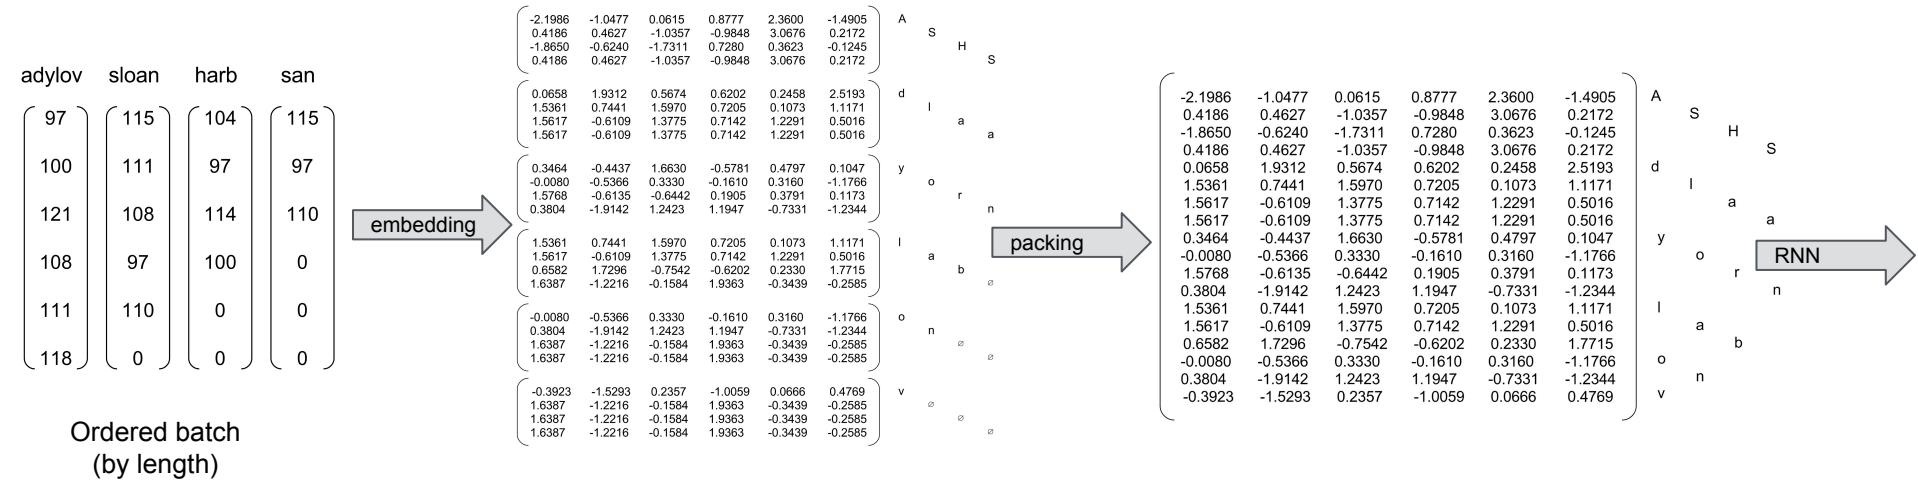
```
classifier = RNNClassifier(N_CHARS, HIDDEN_SIZE, ...
if torch.cuda.is_available():
    classifier.cuda()
```

3

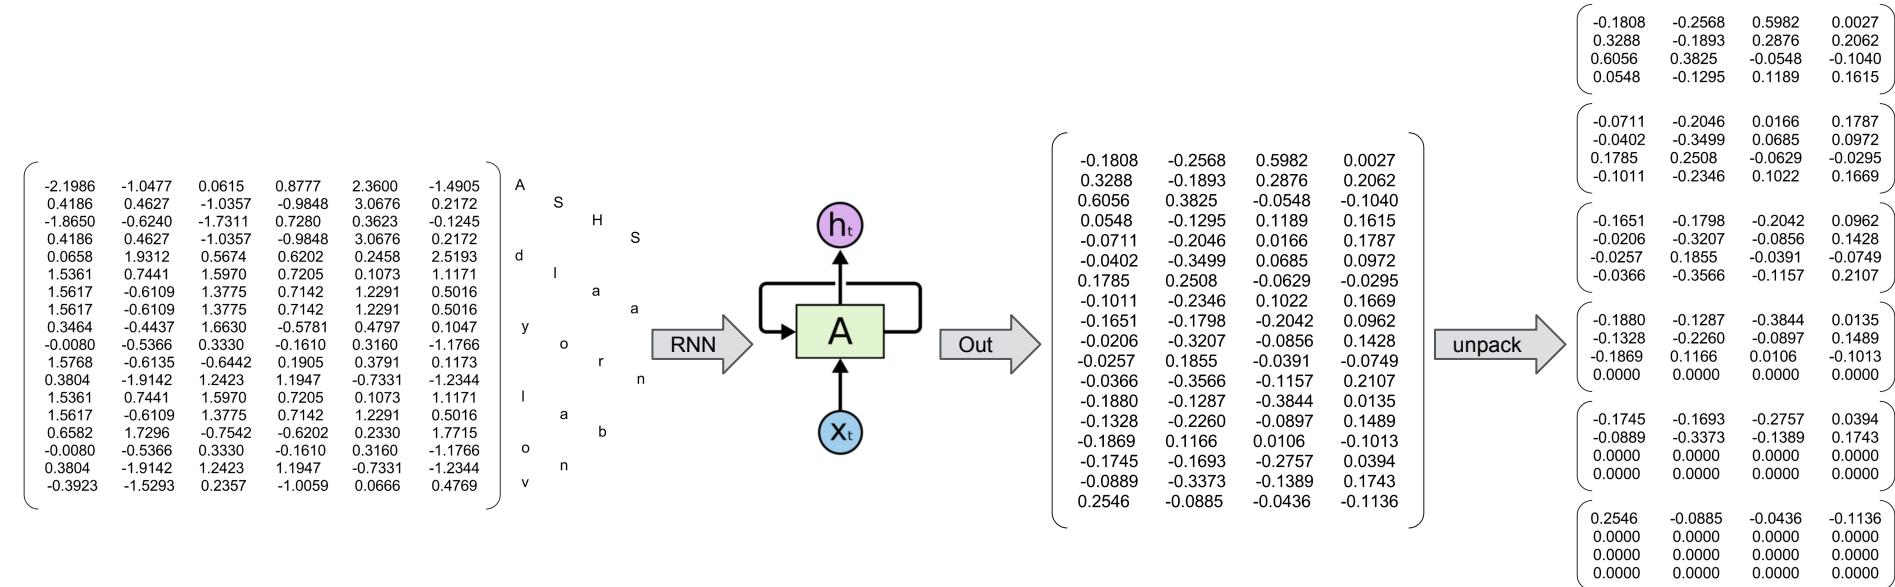
Wrap your model using data parallel

```
if torch.cuda.device_count() > 1:
    print("Let's use", torch.cuda.device_count(), "GPUs!")
    # dim = 0 [33, xxx] -> [11, ...], [11, ...], [11, ...] on 3 GPUs
    classifier = nn.DataParallel(classifier)
```

Efficiently handling batched sequences with variable lengths: pack_padded_sequence



Efficiently handling batched sequences with variable lengths: pack_padded_sequence

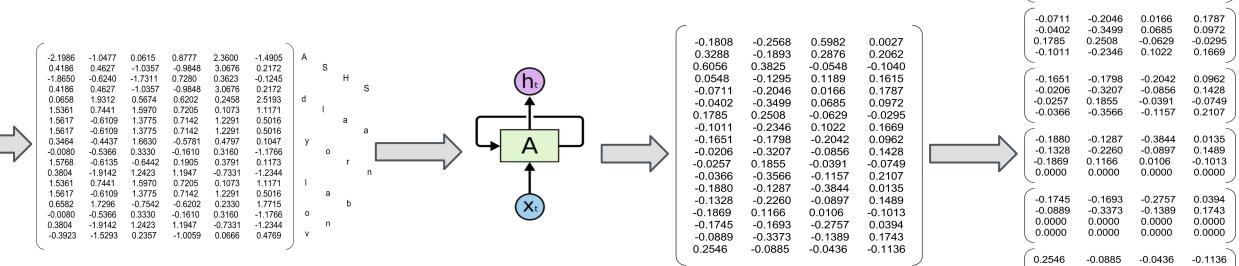


Efficiently handling batched sequences with variable lengths: pack_padded_sequence

	adlyov	sloan	harb	san
97		115	104	115
100	111	97	97	
121	108	114		110
108	97	100	0	
111	110	0	0	
118	0	0	0	0

-2.1996	0.0477	0.0615	0.0777	0.2600	1.4905
0.4196	0.4627	-1.0357	-0.9448	0.0076	0.2172
0.4196	0.4627	-1.0357	-0.9448	0.0076	0.2172
0.4196	0.4627	-1.0357	-0.9448	0.0076	0.2172
0.4196	0.4627	-1.0357	-0.9448	0.0076	0.2172

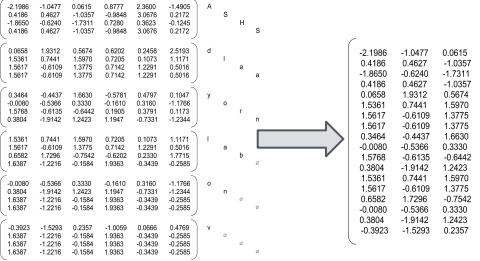
(a) With packing and unpacking



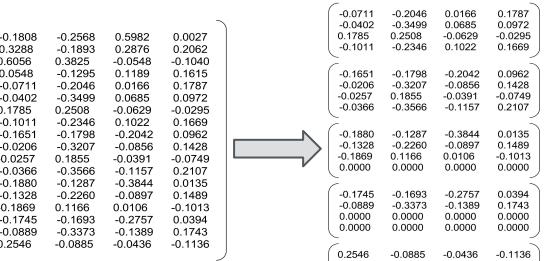
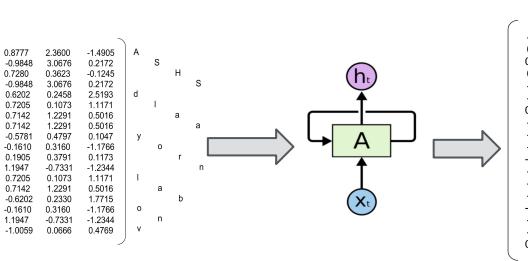
Efficiently handling batched sequences with variable lengths: pack_padded_sequence

adylov Sloan harb san

97	115	104	115
100	111	97	97
121	108	114	110
108	97	100	0
111	110	0	0
118	0	0	0

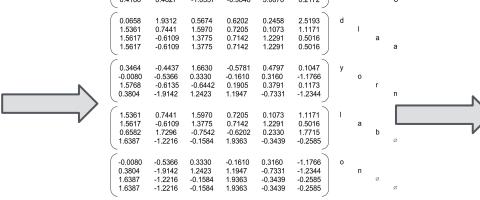


(a) With packing and unpacking

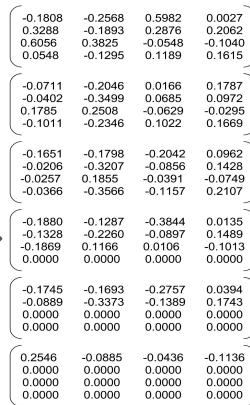


adylov Sloan harb san

97	115	104	115
100	111	97	97
121	108	114	110
108	97	100	0
111	110	0	0
118	0	0	0



(b) Without packing and unpacking



Matrix visualization from Nicolas, <https://github.com/ngarneauu>

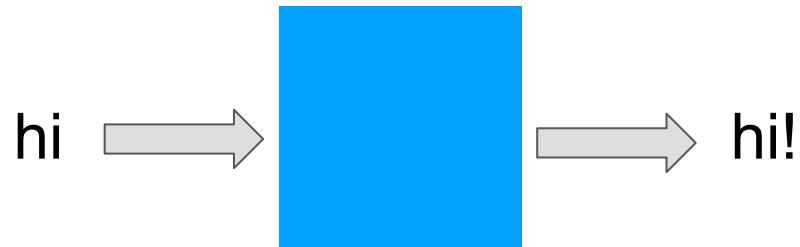
Exercise 13-1: Implement the full name classification

- With GPU
- With data parallel
- Use pad-pack

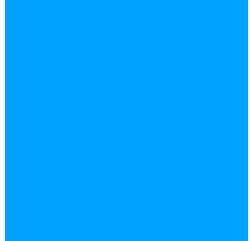
(Character level) Language Modeling



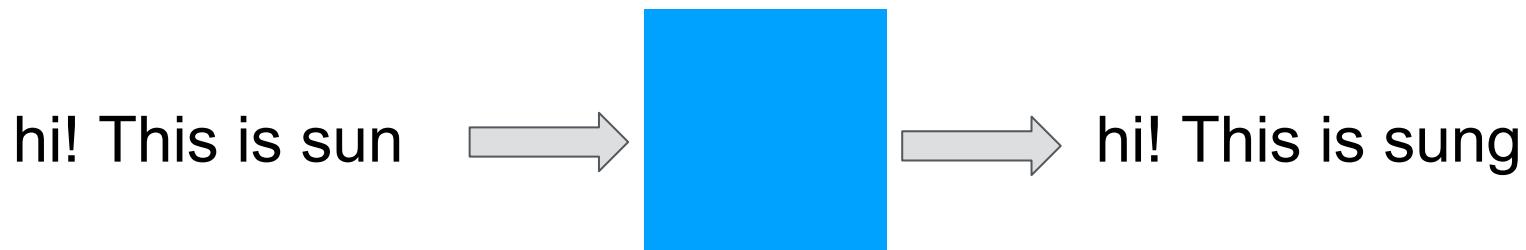
(Character level) Language Modeling



(Character level) Language Modeling

hi! This is sun →  → hi! This is sung

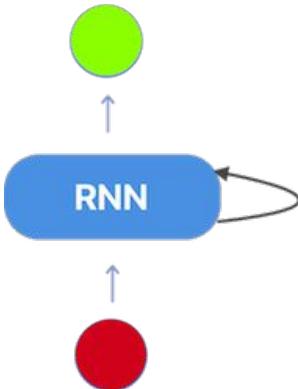
(Character level) Language Modeling



$$p(y_t | y_1, y_2, \dots, y_{t-1})$$

Language Modeling using RNN

i! This is sung



hi! This is sun

```
def generate(decoder, prime_str='A', predict_len=100, temperature=0.8):
    hidden = decoder.init_hidden()
    prime_input = str2tensor(prime_str)
    predicted = prime_str

    # Use priming string to "build up" hidden state
    for p in range(len(prime_str) - 1):
        _, hidden = decoder(prime_input[p], hidden)

    inp = prime_input[-1]
    for p in range(predict_len):
        output, hidden = decoder(inp, hidden)

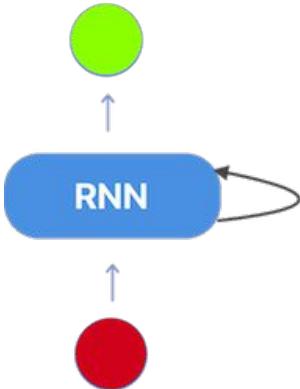
        # Sample from the network as a multinomial distribution
        output_dist = output.data.view(-1).div(temperature).exp()
        top_i = torch.multinomial(output_dist, 1)[0]

        # Add predicted character to string and use as next input
        predicted_char = chr(top_i)
        predicted += predicted_char
        inp = str2tensor(predicted_char)

    return predicted
```

Training

i! This is sung



hi! This is sun

```
def train_teacher_forching(line):
    input = str2tensor(line[:-1])
    target = str2tensor(line[1:])

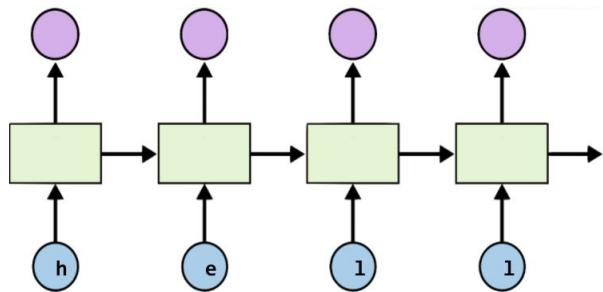
    hidden = RNNdecoder.init_hidden()
    loss = 0

    for c in range(len(input)):
        output, hidden = decoder(input[c], hidden)
        loss += criterion(output, target[c])

    decoder.zero_grad()
    loss.backward()
    decoder_optimizer.step()

    return loss.data[0] / len(input)
```

Teacher Forcing



Teacher Forcing

```
def train_teacher_forcing(line):
    input = str2tensor(line[:-1])
    target = str2tensor(line[1:])

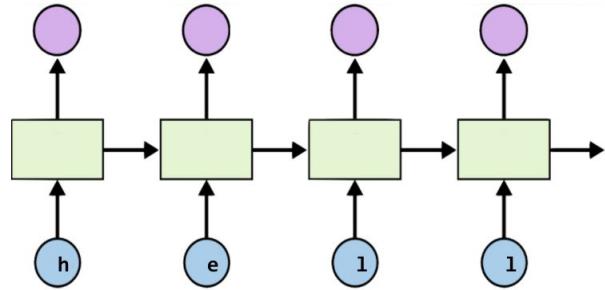
    hidden = RNNdecoder.init_hidden()
    loss = 0

    for c in range(len(input)):
        output, hidden = decoder(input[c], hidden)
        loss += criterion(output, target[c])

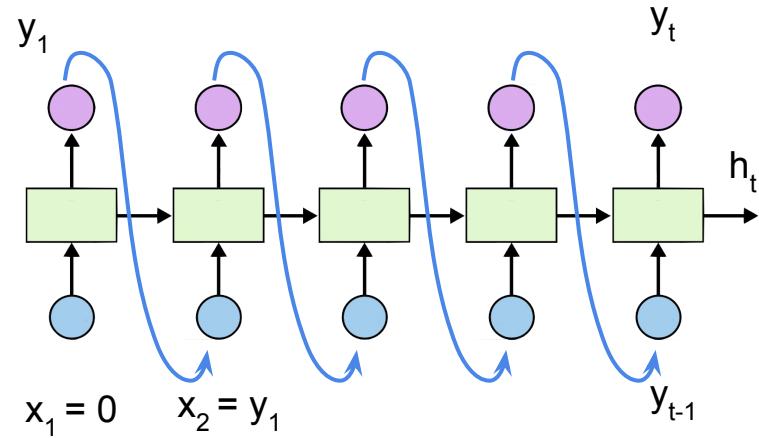
    decoder.zero_grad()
    loss.backward()
    decoder_optimizer.step()

    return loss.data[0] / len(input)
```

Teacher Forcing

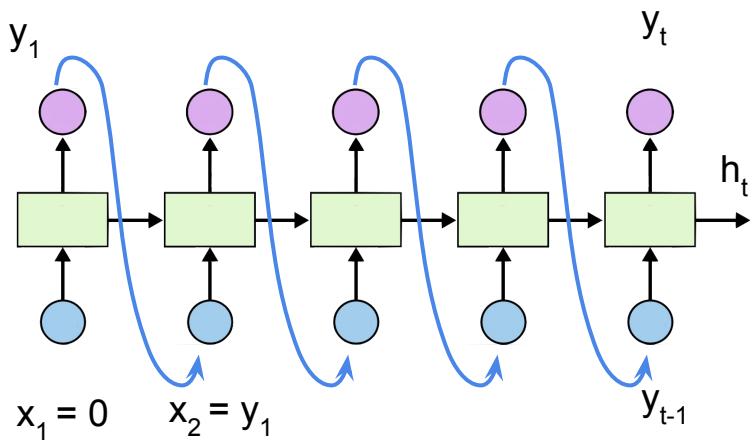


Teacher Forcing



No Teacher Forcing
(more natural)

Training: no teacher forcing



```
def train(line):
    input = str2tensor(line[:-1])
    target = str2tensor(line[1:])

    hidden = decoder.init_hidden()
    decoder_in = input[0]
    loss = 0

    for c in range(len(input)):
        output, hidden = decoder(decoder_in, hidden)
        loss += criterion(output, target[c])
        decoder_in = output.max(1)[1]

    decoder.zero_grad()
    loss.backward()
    decoder_optimizer.step()

    return loss.data[0] / len(input)
```

Exercise 13-2: Language model with Obama speech



Best Speeches of
Barack Obama
through his 2009
Inauguration

**Most Recent Speeches are
Listed First**

- Barack Obama -
Inaugural Speech
- Barack Obama -
Election Night Victory /
Presidential Acceptance Speech -
Nov 4 2008
- Barack Obama - Night Before the
Election - the Last Rally -
Manassas Virginia - Nov 3 2008
- Barack Obama - Democratic
Nominee Acceptance Speech

Obama Inaugural Address 20th January 2009

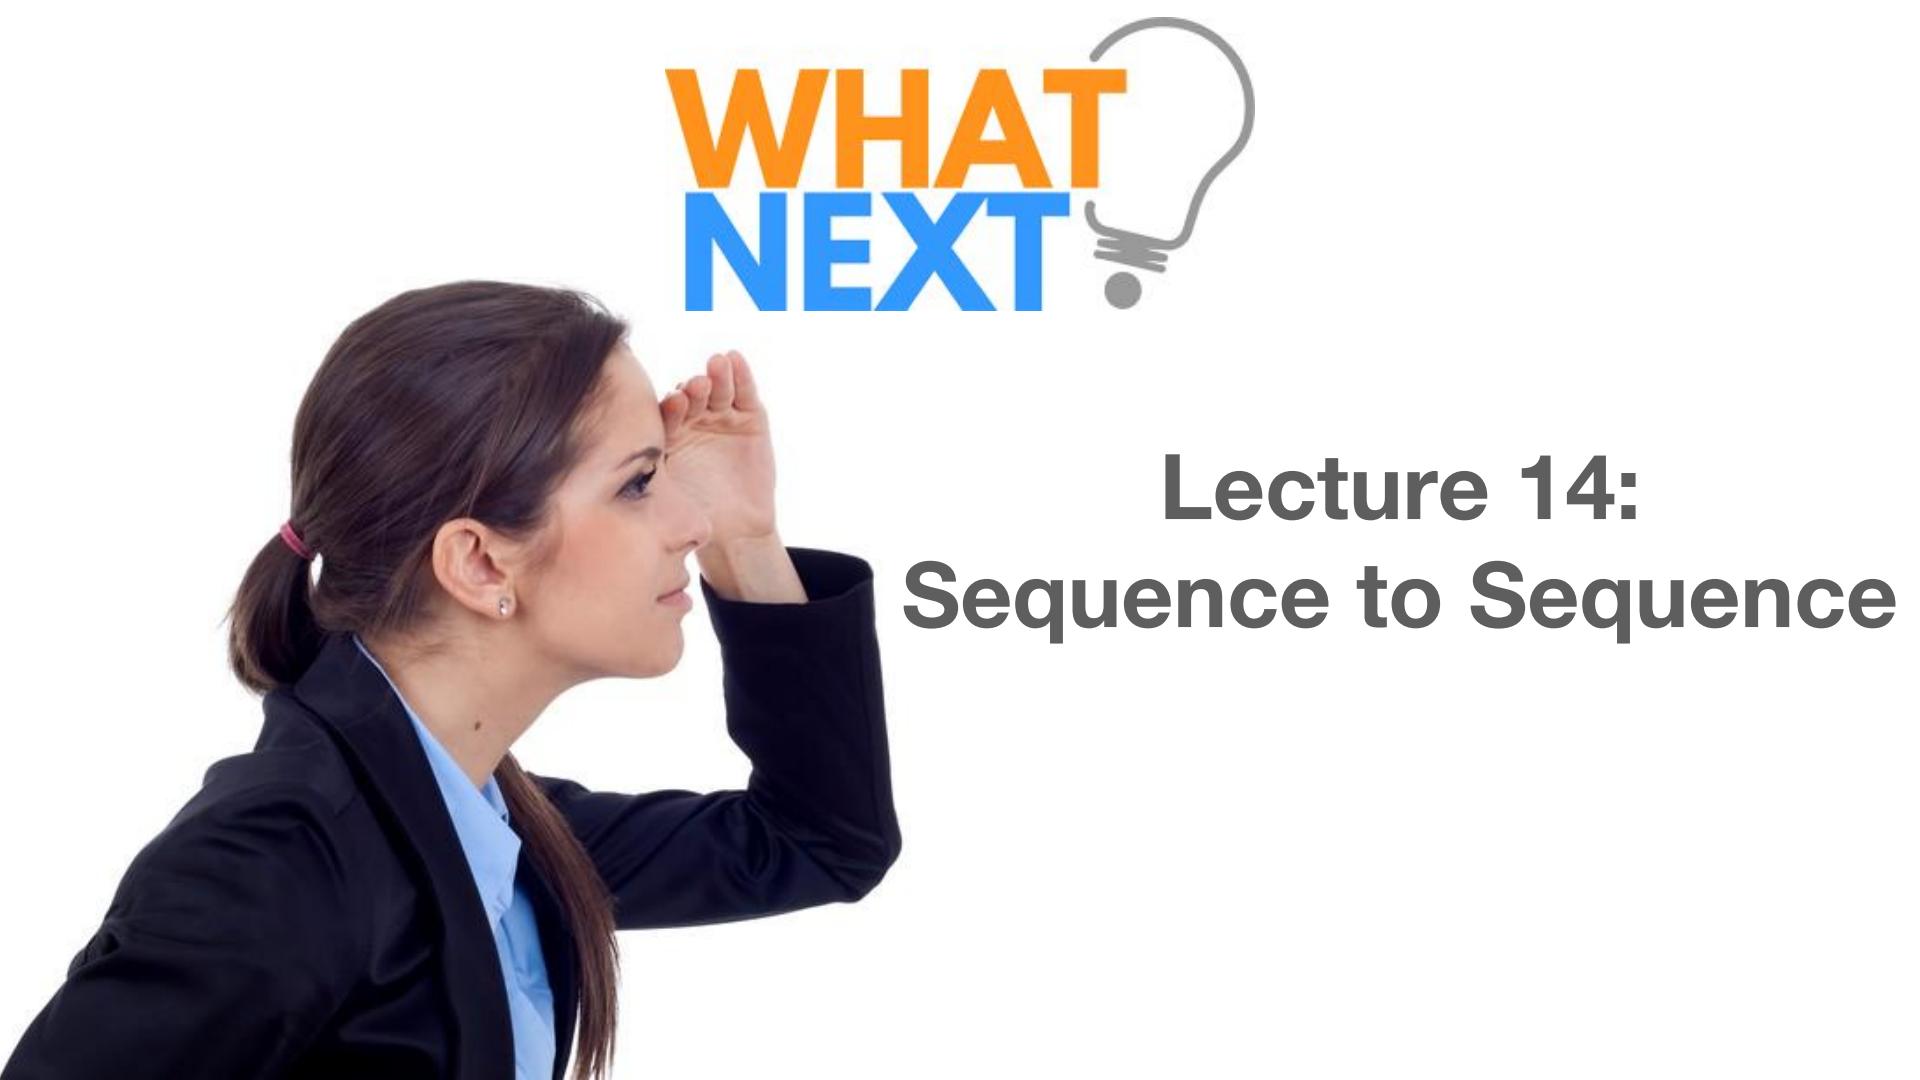
My fellow citizens:

I stand here today humbled by the task before us, grateful for the trust you have bestowed, mindful of the sacrifices borne by our ancestors. I thank President Bush for his service to our nation, as well as the generosity and cooperation he has shown throughout this transition.

Forty-four Americans have now taken the presidential oath. The words have been spoken during rising tides of prosperity and the still waters of peace. Yet, every so often the oath is taken amidst gathering clouds and raging storms. At these moments, America has carried on not simply because of the skill or vision of those in high office, but because We the People have remained faithful to the ideals of our forbearers, and true to our founding documents.

So it has been. So it must be with this generation of Americans.

That we are in the midst of crisis is now well understood. Our nation is at war, against a far-reaching network of violence and hatred. Our economy is badly weakened, a consequence of greed and irresponsibility on the part of some, but also our collective failure to make hard choices and prepare the nation for a new age. Homes have been lost; jobs shed; businesses shuttered. Our health care is too costly; our schools fail too many; and each day brings further evidence that the ways we use energy strengthen our adversaries and threaten our planet.



**WHAT
NEXT**

A woman with long brown hair tied back in a ponytail, wearing a dark blue blazer over a light blue shirt, is shown in profile facing right. She has her right hand raised to her forehead, with her fingers partially covering her eyes as if peering into the distance or shielding her eyes from bright light. Her gaze is directed upwards and to the right. In the upper right corner of the image, there is a graphic element consisting of the words "WHAT NEXT" in large, bold, sans-serif font. The word "WHAT" is in orange and "NEXT" is in blue. To the right of the text is a simple line drawing of a lit lightbulb with a curved wire and a small dot representing the filament.

Lecture 14: Sequence to Sequence